

# Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import sklearn
import matplotlib.pyplot as plt
import seaborn as sns
import time
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: dataSet = pd.read_csv('data/nUsersCreditCardTx.csv')
```

```
In [3]: dataSet.head()
```

```
Out[3]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.01
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.22
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.24
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.10
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.00

5 rows × 31 columns

Notice that the variables are PCA transformed. Due to confidentiality issues the data set available online isn't having the actual feature names. Only Time and Amount are readable ones. While Amount is the total transaction amount, the Time column is the number of seconds elapsed between the transaction in consideration and the very first transaction of the dataset.

```
In [4]: dataSet.info(verbose=True, show_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Time    284807 non-null  float64
 1   V1       284807 non-null  float64
 2   V2       284807 non-null  float64
 3   V3       284807 non-null  float64
 4   V4       284807 non-null  float64
 5   V5       284807 non-null  float64
 6   V6       284807 non-null  float64
 7   V7       284807 non-null  float64
 8   V8       284807 non-null  float64
 9   V9       284807 non-null  float64
10  V10      284807 non-null  float64
11  V11      284807 non-null  float64
12  V12      284807 non-null  float64
13  V13      284807 non-null  float64
14  V14      284807 non-null  float64
15  V15      284807 non-null  float64
16  V16      284807 non-null  float64
17  V17      284807 non-null  float64
18  V18      284807 non-null  float64
19  V19      284807 non-null  float64
20  V20      284807 non-null  float64
21  V21      284807 non-null  float64
22  V22      284807 non-null  float64
23  V23      284807 non-null  float64
24  V24      284807 non-null  float64
25  V25      284807 non-null  float64
26  V26      284807 non-null  float64
27  V27      284807 non-null  float64
28  V28      284807 non-null  float64
29  Amount   284807 non-null  float64
30  Class    284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

Data is clean as there aren't any null values

## Data set limitation

```
In [5]: dataSet['Class'].value_counts()
```

```
Out[5]: 0    284315
        1      492
        Name: Class, dtype: int64
```

Class with value 1 indicate fraudulent transaction and Class with value 0 indicate a legitimate transaction. The data set has only 492 fraudulent transaction which is 0.17% of the data set. This is a very imbalanced data set.

## Data Visualization

### Histogram

```
In [6]: classValues = dataSet['Class'].value_counts().index

fig, (timeClass0Fig, timeClass1Fig) = plt.subplots(2, 1, sharex=True, figsize=(15, 10))

timeClass0Fig.hist(dataSet['Time'][dataSet['Class']==classValues[0]], bins=50)
timeClass0Fig.set_title('Class with the value = ' + str(classValues[0]))
```


```

timeClass1Fig.hist(dataSet['Time'][dataSet['Class']==classValues[1]], bins=50)
timeClass1Fig.set_title('Class with the value = ' + str(classValues[1]))

plt.xlabel('Time')
plt.ylabel('number of TxS')

plt.show()

```



```

In [7]: fig, (timeClass0Fig, timeClass1Fig) = plt.subplots(2, 1, sharex=True, figsize=(15, 10))


timeClass0Fig.hist(dataSet['Amount'][dataSet['Class']==classValues[0]], bins=50)
timeClass0Fig.set_title('Class with the value = ' + str(classValues[0]))

timeClass1Fig.hist(dataSet['Amount'][dataSet['Class']==classValues[1]], bins=50)
timeClass1Fig.set_title('Class with the value = ' + str(classValues[1]))

plt.xlabel('Amount')
plt.ylabel('number of TxS')

plt.show()


```



## Correlation

In [8]:

```
plt.figure(figsize=(10,10))
corr = dataSet.corr()['Class'].sort_values().drop('Class')
corr.plot(kind='barh')
plt.title('Correlation between Features')
plt.grid(True)
plt.show()
```



## Separating the data into featured and dependent variable

```
In [9]: df = dataSet[dataSet['Class'] == 1]
df1 = dataSet[dataSet['Class'] == 0]
df1.drop(df1.index[2000:], 0, inplace=True)
nUsersCreditCardTxs = pd.concat([df, df1], axis=0)
```

```
In [10]: nUsersCreditCardTxs.drop(['V11', 'V4', 'V2', 'V21', 'V19', 'V20', 'V8', 'V27', 'V28', 'Amount', 'Time'], 1, inplace=True)

x = nUsersCreditCardTxs.iloc[:, :-1].values
y = nUsersCreditCardTxs.iloc[:, -1].values
```

```
In [11]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, stratify=y, random_state=42)
```

```
In [12]: from sklearn.preprocessing import StandardScaler

standardScaler = StandardScaler()
x_train = standardScaler.fit_transform(x_train)
x_test = standardScaler.transform(x_test)
```

# Model building

## Model 1 - Logistic Regression

```
In [13]: from sklearn.linear_model import LogisticRegression

lrModel = LogisticRegression()
lrModel = lrModel.fit(x_train,y_train)
y_pred = lrModel.predict(x_test)
```

```
In [14]: from sklearn.metrics import accuracy_score, confusion_matrix


accuracyLR = accuracy_score(y_test,y_pred)
confusionMatrixLR = confusion_matrix(y_test,y_pred)
```

```
In [15]: print(f"Accuracy is {accuracyLR*100}")
print("Confusion Matrix is ")
print(confusionMatrixLR)
```

```
Accuracy is 95.99198396793587
Confusion Matrix is
[[399  1]
 [ 19 80]]
```

```
In [16]: from sklearn.metrics import plot_confusion_matrix

plot_confusion_matrix(lrModel, x_test, y_test, cmap=plt.cm.Greens)
plt.show()
```



## Model 2 - Linear SVM (Support Vector Machine)

```
In [17]: from sklearn import svm

svmModel = svm.SVC(kernel='linear')
```

```
In [18]: svmModel.fit(x_train, y_train)
```

```
Out[18]: SVC(kernel='linear')
```


```
In [19]: y_pred = svmModel.predict(x_test)
```

```
In [20]: accuracySVM = accuracy_score(y_test,y_pred)
confusionMatrixSVM = confusion_matrix(y_test,y_pred)
```

```
In [21]: print(f"Accuracy is {accuracySVM*100}")
print("Confusion Matrix is ")
print(confusionMatrixSVM)
```

```
Accuracy is 95.39078156312625
Confusion Matrix is
[[399   1]
 [ 22  77]]
```

```
In [22]: plot_confusion_matrix(svmModel, x_test, y_test, cmap=plt.cm.Blues)
plt.show()
```



## Model 3 - Decision Tree

```
In [23]: from sklearn import tree

dtModel = tree.DecisionTreeClassifier()
dtModel = dtModel.fit(x_train,y_train)
y_pred = dtModel.predict(x_test)
```


```
In [24]: accuracyDt=accuracy_score(y_test,y_pred)
confusionMatrixDt=confusion_matrix(y_test,y_pred)
```

```
In [25]: print(f"Accuracy is {accuracyDt*100}")
print("Confusion Matrix is ")
print(confusionMatrixDt)
```

```
Accuracy is 95.39078156312625
Confusion Matrix is
[[392   8]
 [ 15  84]]
```

```
In [26]: plot_confusion_matrix(dtModel, x_test, y_test, cmap=plt.cm.Reds)
```

```
plt.show()
```



## Model 4 - Random Forest Classifier

```
In [27]: from sklearn.ensemble import RandomForestClassifier


rfModel =RandomForestClassifier()
rfModel = rfModel.fit(x_train,y_train)
y_pred = rfModel.predict(x_test)
```

```
In [28]: accuracyRf=accuracy_score(y_test,y_pred)
confusionMatrixRf=confusion_matrix(y_test,y_pred)
```

```
In [29]: print(f"Accuracy is {accuracyRf*100}")
print("Confusion Matrix is ")
print(confusionMatrixRf)
```

```
Accuracy is 96.59318637274549
Confusion Matrix is
[[399  1]
 [ 16  83]]
```

```
In [30]: plot_confusion_matrix(rfModel, x_test, y_test,cmap=plt.cm.Greens)
plt.show()
```





# Model 5 - VotingClassifier

```
In [31]: from sklearn.ensemble import VotingClassifier


vcModel = VotingClassifier(estimators=[('rf', rfModel), ('dt', dtModel)], voting='hard')
vcModel = vcModel.fit(x_train,y_train)
y_pred = vcModel.predict(x_test)
```

```
In [32]: accuracyVC = accuracy_score(y_test,y_pred)
confusionMatrixVC = confusion_matrix(y_test,y_pred)
```

```
In [33]: print(f"Accuracy is {accuracyVC*100}")
print("Confusion Matrix is ")
print(confusionMatrixVC)
```

```
Accuracy is 96.39278557114228
Confusion Matrix is
[[399   1]
 [ 17  82]]
```

```
In [34]: plot_confusion_matrix(vcModel, x_test, y_test,cmap=plt.cm.Blues)
plt.show()
```



# Model 6 - XGBoost

```
In [35]: import xgboost as xgb

#xgbModel = xgb.XGBRegressor(objective='reg:logistic', colsample_bytree = 0.3, learning_rate = 0.1,
#                             max_depth = 5, alpha = 10, n_estimators = 10)
xgbModel = xgb.XGBClassifier()
xgbModel = xgbModel.fit(x_train,y_train)
y_pred = xgbModel.predict(x_test)
```


```
[15:52:56] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
```

```
In [36]: accuracyXGB = accuracy_score(y_test,y_pred)
confusionMatrixXGB = confusion_matrix(y_test,y_pred)
```

```
In [37]: print(f"Accuracy is {accuracyXGB*100}")
print("Confusion Matrix is ")
print(confusionMatrixXGB)
```

```
Accuracy is 96.39278557114228
Confusion Matrix is
[[398   2]
 [ 16  83]]
```

```
In [38]: plot_confusion_matrix(xgbModel, x_test, y_test, cmap=plt.cm.Reds)
plt.show()
```



```
In [39]: print(f"Model 1 - Logistic Regression model accuracy is {accuracyLR*100}")
print(f"Model 2 - Linear SVM model accuracy is {accuracySVM*100}")
print(f"Model 3 - Decision tree model accuracy is {accuracyDt*100}")
print(f"Model 4 - Random forest classifier model accuracy is {accuracyRf*100}")
print(f"Model 5 - Voting classifier model accuracy is {accuracyVC*100}")
print(f"Model 6 - XGBoost model accuracy is {accuracyXGB*100}")
```

```
Model 1 - Logistic Regression model accuracy is 95.99198396793587
Model 2 - Linear SVM model accuracy is 95.39078156312625
Model 3 - Decision tree model accuracy is 95.39078156312625
Model 4 - Random forest classifier model accuracy is 96.59318637274549
Model 5 - Voting classifier model accuracy is 96.39278557114228
Model 6 - XGBoost model accuracy is 96.39278557114228
```

Given the data set, comparing accuracy of the models from above, the suggested model is Random forest classifier or XGBoost. For simplicity Random forest classifier model has been saved for later execution.

```
In [40]: import joblib

joblib.dump(rfModel, "./random_forest.joblib")

#Loading the model at a later execution stage
#loaded_rfModel = joblib.load("./random_forest.joblib")
```

```
Out[40]: ['./random_forest.joblib']
```